

Week 3 Questions

Karpathy GPT

1. In the forward function of the GPT class, why is there not a call written to apply the softmax function to the logits?
2. What does `torch.no_grad()` on line 282 in `model.py` do and why is it a good idea to use this?
3. Why do they transpose the tensors on lines 57-59 in `model.py`?
4. What does the `self.register_buffer` call on line 47 of `model.py` do? What would go wrong if instead we added the tensor in this line to the model parameters?
5. Try training and evaluating the transformer on some data of your choice. You may want to use Google Colab, which gives gpu access for free, to run the code. Note that using a gpu will make training and evaluation go much more quickly than cpu.

Language Models are Few-Shot Learners

1. What do the authors mean by 'zero-shot', 'one-shot', and 'few-shot'?
2. In what ways is few-shot learning different from fine-tuning?
3. They use a filtered version of Common crawl for most of their training data. What advantages and disadvantages are there for using this data?
4. The mention in section 2.4 that larger numbers of in-context examples are not always better. Why might this be the case?
5. In section 3.9.4, they show that news articles generated by the 175B version of GPT-3 are nearly indistinguishable from real news articles. What limitations might there be in this section?
6. In section 6.2 the authors do a short analysis on some of the biases GPT-3 has. What are some reasons the model has these biases?