

Week 8 Questions

Large Language Models are Zero-Shot Reasoners

Brief Summary: They find substantially improved performance on multiple benchmarks with a new prompting method. Their method is simple, they ask the model to think step-by-step before determining the answer to a question.

1. Describe their method.
2. Compare/contrast the 4 prompting methods in figure 1 of the paper.
3. Compare their method to scratchpad fine-tuning from a few classes ago.

Generative Agents: Interactive Simulacra of Human Behavior

Brief Summary: The paper uses a language model to generate actions and dialogue for agents in a game. Memories of past observations are stored for each agent and recalled to inform their current actions and dialogue. The actions and dialogue for each agent are generated by prompting a language model. Multiple agents interact in the environment.

1. They provide some dialogue examples from the agents. What do you think about them?
2. Describe some limitations of their method and simulation.
3. How are memories stored and retrieved?
4. How does the planning work?